When Does the Machine Judge Fail?
An Empirical Study of LLM-as-a-Judge for English, Hindi Machine Translation

Ekansh Goyal
2023114007
ekansh.goyal@research.iiit.ac.in

Abstract

The LLM-as-a-Judge paradigm offers a scalable alternative to human evaluation, yet its reliability for morphologically rich target languages remains poorly understood. We present a systematic empirical study using gemini 3.1 Flash lite as a judge for English-to-Hindi machine translation, comparing Qwen3-32B and GPT-OSS-20B across 50 curated examples from FLORES-200. We examine three evaluation protocols (direct scoring, pairwise comparison, rubric-based), a four-variant prompt sensitivity study, and a 12-example adversarial test suite. Our principal findings are: (i) rubric-based evaluation achieves the highest human, judge correlation ($r_s = 0.460$, $p < 0.001$), while BLEU fails to reach significance; (ii) the judge is not deceived by any adversarial manipulation, contrary to prevailing assumptions; (iii) chain-of-thought prompting degrades rather than improves judgment quality; and (iv) human inter-annotator agreement is itself very low ($\bar{\kappa} = 0.132$) on challenging inputs, which materially complicates the interpretation of moderate LLM, human correlations. [1]

1 Introduction

Evaluating machine translation at scale is expensive: human annotation is slow, subject to inconsistency, and difficult to reproduce. The LLM-as-a-Judge paradigm (Zheng et al., 2023; Liu et al., 2023; Kocmi and Federmann, 2023) proposes using large language models as surrogate annotators, promising automation without a substantial quality penalty. Yet the conditions under which such judges are reliable, and those under which they fail, remain poorly characterized, particularly for morphologically rich, non-Latin-script language pairs.

This study treats the judge itself as the object of inquiry. We choose English-to-Hindi translation because Hindi presents systematic structural challenges: subject-object-verb order, postpositions, grammatical gender agreement, and Devanagari script all create distinctive error modes for both MT systems and automated evaluators. We deliberately construct a benchmark that covers five distinct challenge types known to stress-test automated evaluation, then examine where the judge succeeds, where it fails, and what kind of analysis is required to tell the difference.

Our contributions are as follows. We release a five-category benchmark of 50 examples from FLORES-200 designed to expose specific LLM evaluator weaknesses. We systematically compare three evaluation protocols against human gold with all significance levels reported. We present a 12-example adversarial suite revealing unexpected judge robustness. Finally, we develop an empirically grounded failure mode taxonomy and use it to argue that the dominant failure mode in this setting is *not* the one most commonly assumed.

2 Experimental Setup

2.1 MT Systems

We evaluate two instruction-tuned LLMs in zero-shot translation mode, selected to represent architecturally distinct model families from different organizations.

System A (Qwen3-32B) (Qwen Team, 2025) is a 32-billion-parameter model by Alibaba Cloud with strong multilingual capabilities. **System B (GPT-OSS-20B)** (OpenAI, 2025) is a 20-billion-parameter open-weights model by OpenAI. Both are accessed via the Groq inference API. Selecting systems from different organizations avoids intra-family confounds in the quality comparison, and choosing two substantially different model sizes tests whether scale alone predicts translation quality differences.

[1] Code, datasets, and prompts are publicly available at: https://github.com/Ekansh0301/EMNLP_A4

2.2 Dataset

We draw from the FLORES-200 (NLLB Team et al., 2022) English-Hindi devtest split, which provides professionally human-translated sentence pairs sourced from Wikipedia, Wikinews, and WikiVoyage. These translations serve as evaluation references without modification.

To maximize coverage of phenomena that are known to challenge LLM evaluators, we select 50 examples across five categories of 10 examples each, identified through deterministic filters over the full devtest set:

Ambiguous. Sentences containing lexical polysemy, pronoun reference ambiguity, or scope ambiguity that admits multiple valid Hindi renderings, making any single judgment contestable.

Hallucination-prone. Sentences dense with named entities, specific dates, percentages, and numerical values that MT systems frequently confabulate or transpose.

Noisy inputs. Clean FLORES sentences with programmatically injected real-world noise: keyboard typos, internet slang, English-Hindi code-switching (Hinglish), OCR character corruption, and missing punctuation. The original FLORES Hindi translation is retained as the reference, testing whether MT systems handle input noise gracefully.

Negation/quantifier traps. Sentences where dropped or altered scope-bearing elements entirely reverse meaning, modal verbs (*could have* vs. *did*), frequency quantifiers (*rarely* vs. *sometimes*), and scalar expressions (*almost all* vs. *all*).

Long and complex. Sentences exceeding 22 words with subordinate clauses, relative clauses, or multiple comma-separated constituents, requiring substantial English-SVO to Hindi-SOV reordering.

3 Human Evaluation

3.1 Annotation Protocol

Three annotators with native or near-native Hindi proficiency evaluated all 50 examples via a structured Google Form. For each example, annotators viewed the English source, the FLORES reference, and both MT outputs side by side. Two dimensions were rated independently on 1, 5 Likert scales: **fluency** (grammaticality and naturalness of the Hindi) and **adequacy** (completeness of meaning transfer from English). Annotators also provided an overall preference judgment (System A / System B / Tie). A rubric card with scoring anchors was distributed before annotation. For noisy-input examples, annotators were instructed to evaluate based on the intended clean meaning of the source.

3.2 Inter-Annotator Agreement

Measure	Dimension	Score
Krippendorff α (ordinal)	Sys A, Adequacy	-0.116
	Sys A, Fluency	-0.123
	Sys B, Adequacy	-0.129
	Sys B, Fluency	-0.060
	Mean	-0.107
Cohen κ (preference)	Ann. 1 vs. Ann. 2	0.125
	Ann. 1 vs. Ann. 3	0.140
	Ann. 2 vs. Ann. 3	0.131
	Mean	0.132

Table 1: Inter-annotator agreement. Negative Krippendorff α indicates systematic disagreement beyond chance. Cohen $\kappa = 0.132$ corresponds to slight agreement (Landis and Koch, 1977).

Agreement is very low across all measures (Table 1). Negative Krippendorff α values arise when annotators disagree more than would be expected by chance, typically because each rater uses the scale in a systematically different direction rather than randomly. Standard reliability thresholds treat $\alpha < 0.4$ as unacceptable (Krippendorff, 2011); all four of our Likert dimensions fall well below this threshold.

Critically, this finding contextualizes all subsequent LLM, human correlation results. A moderate judge, human correlation ($r_s \approx 0.44$) measured against an unreliable gold standard may represent a ceiling effect rather than a gap to close. We retain the aggregated mean annotations as a best-available proxy, while acknowledging that instability in the gold is a confound throughout the analysis.

3.3 Gold Standard Construction

Gold scores are computed as the mean of three annotators per (example, system, dimension) cell. Preference is resolved by majority vote. System B is consistently rated higher: mean overall score 3.963 vs. 3.803 for System A. It is preferred in 52% of examples (34% System A, 14% Tie). The quality gap is modest but consistent across both fluency and adequacy dimensions.

4 LLM as Judge

We use **gemini 3.1 Flash lite** as the primary judge throughout all experiments. This model is from Google DeepMind, architecturally and organizationally distinct from both MT systems, avoiding intra-family preference confounds. All judge calls use temperature 0.1. Full prompts are listed in Appendix B.

4.1 Direct Scoring

Each (source, translation, reference) triple is presented to the judge with a 1, 5 rating scale and an instruction to return a single integer. The judge assigns mean scores of 3.28 and 4.06 to Systems A and B respectively, a gap of 0.78 points. This amplifies the human quality gap (0.16 points) by nearly a factor of five, indicating that the judge is more decisive than humans about which system is better.

4.2 Pairwise Comparison

Each translation pair is judged twice: once in the original order and once with candidates swapped. Swapped verdicts are inverted back to original system labels, enabling direct measurement of position bias. The judge prefers System B in the majority of comparisons, consistent with direct scoring results.

4.3 Rubric-Based Evaluation

The judge scores four criteria via a JSON-format response: adequacy, fluency, terminology handling, and style match. System B receives higher scores on all four dimensions, with the largest gap in terminology (3.78 vs. 4.60), suggesting sensitivity to how well each system preserves proper names, numbers, and domain-specific expressions. Among the three protocols, rubric-based evaluation achieves the highest correlation with human gold (Section 5).

4.4 Prompt Sensitivity

Prompt Variant	Mean	Std.
Verbose expert persona	2.82	1.11
Standard with reference	3.20	0.99
No reference	3.22	1.01
Minimal (one-line)	3.30	1.12

Table 2: Prompt sensitivity for System A ($n = 50$ per variant). Persona framing causes the largest shift; reference omission has negligible effect.

We evaluate System A under four prompt variants varying along two dimensions: wording formality and reference presence (Table 2). Two results are noteworthy. First, the standard-with-reference and no-reference variants differ by only 0.02 points on average, the judge does not anchor its scoring on the reference as a quality ceiling, which is a positive reliability property. Second, framing the judge as a "certified linguist with 20 years of experience" shifts mean scores downward by 0.38, 0.48 points relative to all other variants. The judge's calibration responds to implied expertise level: a more expert persona applies a stricter implicit threshold even when the evaluation criteria are identical.

5 Meta-Evaluation

Metric Pair	r_s	p	τ	p
Direct vs. Human	0.440	<0.001	0.367	<0.001
Rubric vs. Human	0.460	<0.001	0.356	<0.001
Rub. Fluency vs. Human	0.395	<0.001	0.324	<0.001
Rub. Adequacy vs. Human	0.436	<0.001	0.381	<0.001
BLEU vs. Human	0.153	0.128	0.118	0.097

Table 3: Spearman (r_s) and Kendall (τ) correlations between automated metrics and human gold ($n = 100$: 50 examples $\times$ 2 systems). All LLM metrics significant; BLEU is not.

Human, judge correlation. All LLM-based metrics correlate significantly with human gold, while BLEU ($r_s = 0.153$, $p = 0.128$) fails to reach significance (Table 3, Figure 1). The rubric-based overall score performs best, and this advantage is consistent across both Spearman and Kendall measures. BLEU's failure is expected for morphologically rich target languages: correct Hindi translations can use different but valid inflected forms against a single reference, receiving low n-gram overlap scores regardless of quality.

The moderate range of $r_s = 0.395, 0.460$ for LLM metrics should be interpreted carefully. Given mean $\kappa = 0.132$ across annotator pairs, the gold standard is itself noisy. Moderate LLM, human correlation may therefore reflect shared difficulty rather than systematic judge error.

Position bias. Pairwise order consistency is 86% overall (Figure 1, bottom center). The 14% of inconsistent judgments are not uniformly distributed: the long/complex category shows lower consistency, consistent with the hypothesis that lengthy,

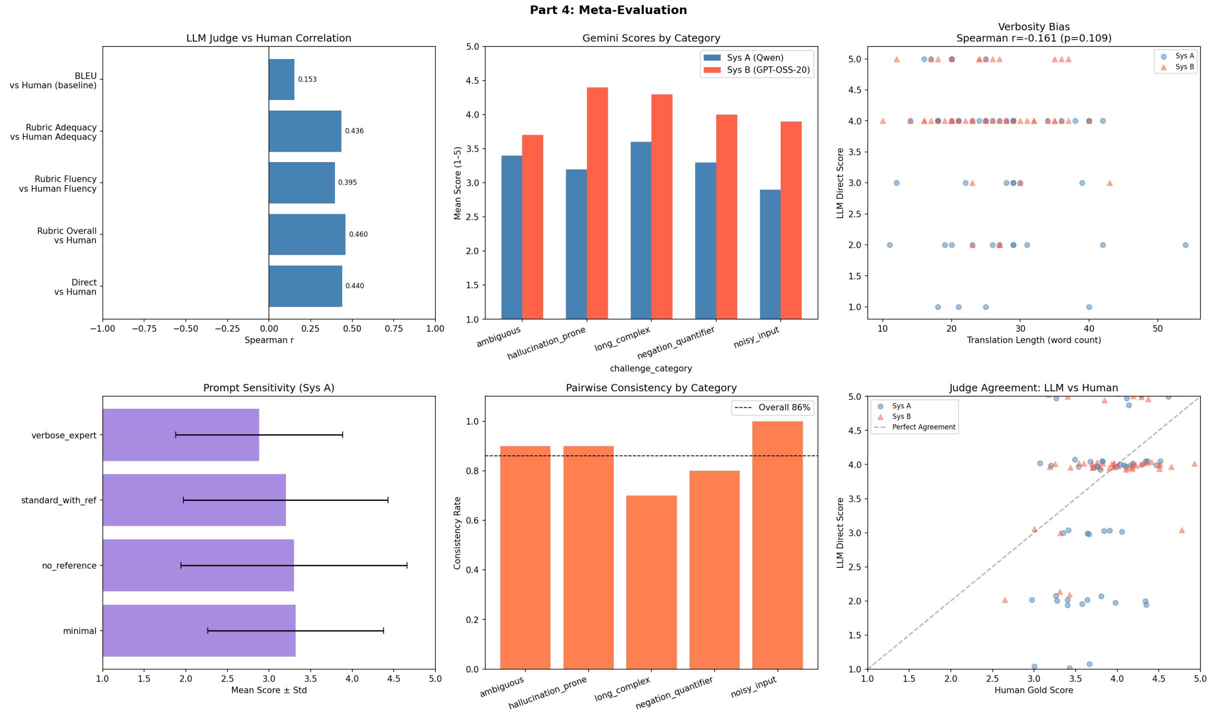

Figure 1: Meta-evaluation overview. Top row: Spearman correlations vs. human gold; mean LLM scores by challenge category and system; verbosity bias scatter. Bottom row: prompt sensitivity distributions; pairwise consistency by category; direct scores vs. human gold scatter.

structurally dense translations more easily capture the judge's attention when presented first.

Verbosity bias. Spearman correlation between translation word count and direct score is $r_s = -0.161$, $p = 0.109$, not significant. No evidence of verbosity bias is present; if anything, the slightly negative sign suggests marginal preference for concise outputs, but this does not reach significance.

Reference bias. A paired t-test comparing the standard-with-reference and no-reference variants reveals no significant difference. The judge evaluates translation quality relative to the source rather than reference proximity, which supports its use in settings where multiple valid reference translations exist.

Self-preference bias. Self-preference analysis requires a judge that generated at least one of the MT system outputs. Since the judge (Gemini) was used exclusively for evaluation and produced no translations in this study, classical self-preference bias analysis is not applicable. This is by design: using a judge from a different model family than both MT systems eliminates this confound entirely.

Type	n	Adv.	Ref.	Fooled
Fluent-but-wrong	4	2.00	5.00	0/4
Correct-but-poor	4	1.50	5.00	0/4
Meaning shift	4	2.25	5.00	0/4
Overall	12	1.92	5.00	0/12

Table 4: Adversarial test results. Mean scores for adversarial and reference translations; the judge correctly preferred the reference in all 12 cases.

6 Adversarial Testing

We construct 12 adversarial translation pairs covering three manipulation types designed to probe the vulnerabilities most commonly attributed to LLM judges.

Fluent-but-wrong translations are grammatically natural Hindi with a single semantic error: directional inversions ("leave hospital" translated as "go to hospital"), factual inversions ("increased" as "decreased"), and subject-object reversals ("received" as "gave").

Correct-but-poor translations preserve the English meaning faithfully but are written in unnatural Hindi with heavy English borrowing, incorrect case markers, and scrambled word order, testing whether the judge rewards correctness even when surface form is poor.

Meaning-shift translations appear as plausible paraphrases but alter truth-conditional content through dropped quantifiers ("almost all" becomes "all"), dropped modality ("could have committed" becomes "committed"), and frequency shifts ("rarely" becomes "sometimes").

The judge was not fooled by a single example across all three types (Table 4). The mean adversarial score of 1.92 against a reference score of 5.00 indicates confident rejection, not borderline hesitation. This result is substantially more favorable than anticipated and runs counter to widely reported claims that LLM judges are susceptible to fluency manipulation (Shen et al., 2023).

Two important caveats apply. First, these adversarial manipulations are relatively coarse, a single identifiable semantic error per example. Subtler perturbations (e.g., coreference errors with no surface evidence, culturally inappropriate register, or pragmatic implicature violations) may not be detected. Second, the judge's success on factual inversions may depend on world-knowledge access; this would not generalize to low-resource domains where such knowledge is absent.

7 Failure Analysis

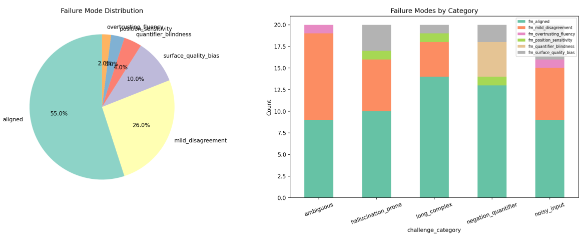

Figure 2: Failure mode distribution across all 100 judgments. Aligned judgments constitute a majority; surface quality bias is the most prevalent non-aligned failure.

We classify each LLM judgment relative to the human gold using a six-category taxonomy (Figure 2). A judgment is *aligned* when the score gap is at most 0.5 points; larger discrepancies are assigned to failure modes based on direction, magnitude, and the challenge category of the example.

Aligned judgments (55%) confirm that the judge agrees with humans on a majority of cases. Agreement is strongest in the long/complex category, where structural quality differences are unambiguous, and weakest in the ambiguous category, where no single score is uniquely defensible.

Mild disagreement (26%) captures moderate discrepancies (0.5, 1.5 points) in either direction without a dominant pattern. This mode is most prevalent for ambiguous examples and reflects genuine uncertainty rather than systematic bias.

Surface quality bias (10%) is the most important failure mode: the judge scores translations at least 1.5 points *below* the human mean, penalizing outputs that are semantically correct but rough in surface form. This directly inverts the failure mode assumed by most LLM evaluation criticism, the judge under-rewards correctness when fluency is poor, rather than over-rewarding fluency when correctness is absent.

Quantifier blindness (4%) is concentrated in the negation/quantifier category. The judge correctly penalizes most meaning shifts, but in a small fraction of cases fails to detect subtle scope alterations that produce semantically incorrect but natural-sounding output.

Position sensitivity (4%) marks cases where the pairwise verdict changes with candidate ordering, concentrated in long/complex examples where longer texts shift the judge's attention.

Overtrusting fluency (2%), the mode most commonly attributed to LLM judges, is the rarest in our data. The judge is less susceptible to fluency manipulation than prior literature suggests.

A notable cross-category pattern: noisy-input examples produce the most diverse failure profile, with all six modes represented. This reflects the absence of a clear evaluation anchor: when the source itself is malformed, neither humans nor LLMs have a stable reference point.

8 Critical Discussion

When does LLM evaluation align with humans? The judge aligns most reliably when quality differences are large and structurally evident, long/complex sentences with clear reordering errors, or hallucination-prone sentences where the judge's world knowledge flags incorrect facts. Alignment is lowest for ambiguous and noisy-input examples. Crucially, these are the same categories where human annotators disagree most with each other (Table 1), which suggests the difficulty is inherent to the evaluation task rather than specific to the judge.

Where does it fail and why? The dominant failure mode is surface quality bias (10

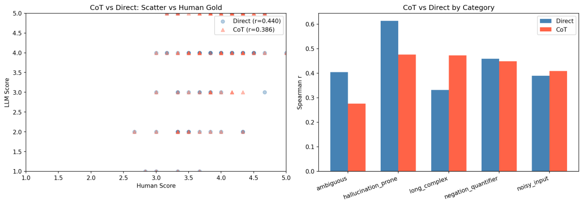

Figure 3: Chain-of-thought vs. direct scoring. Left: LLM scores vs. human gold for both modes. Right: Spearman r_s by challenge category. Direct scoring outperforms CoT overall and in the ambiguous category specifically.

Is it internally consistent? At 86% pairwise order consistency, the judge is substantially but not perfectly consistent. Score range across four prompt variants (approximately 0.48 points, 12% of the full scale) indicates moderate prompt sensitivity. The judge is reliable enough for ranking comparisons between clearly distinct systems; it is not reliable for absolute score comparisons across differently-worded prompts.

How sensitive is it to prompts? Reference presence has negligible effect, the judge evaluates source-translation fidelity rather than reference proximity. Persona framing has the largest effect: a stricter implied expertise level shifts scores downward by nearly 0.5 points. This is a calibration sensitivity, not a validity concern per se, but it means that scores from differently-framed prompts cannot be directly compared.

Can it replace human evaluation? For development-stage system comparison with a clear quality gap, LLM evaluation is reliable enough to serve as a first-pass proxy. For publication-quality research comparing closely matched systems, it cannot substitute for human judgment. The critical caveat from this study is that the human gold standard is itself unstable ($\kappa = 0.132$), which complicates the normative question: if humans disagree at this level, what constitutes the "right" evaluation in the first place?

9 Extension: Chain-of-Thought vs. Direct Scoring

We compare two elicitation strategies: direct scoring (single-integer response) and chain-of-thought (CoT) prompting (four explicit reasoning steps before the final score). Contrary to findings in general-purpose benchmarks, CoT prompting degrades correlation with human gold: $r_s = 0.386$ vs. 0.440 for direct scoring (Figure 3).

The category-level breakdown reveals that this degradation is not uniform. CoT underperforms most severely on ambiguous examples ($r_s = 0.276$ vs. 0.405 for direct scoring), while performing comparably to direct scoring on noisy-input examples (0.410 vs. 0.388). Negation/quantifier and hallucination-prone categories show near-identical performance between the two modes.

Two mechanisms may explain the ambiguous-category degradation. First, *reasoning drift*: when forced to produce explicit intermediate judgments (e.g., "meaning is partially preserved"), the judge may commit to these assessments in ways that resist correction, producing more extreme final scores than holistic assessment would yield. Second, *implicit criterion reweighting*: the four CoT steps treat meaning, terminology, grammar, and register as sequentially equal, whereas human holistic judgments weight these dynamically by context. For ambiguous inputs, committing to a single reading of the ambiguity before scoring compounds the evaluation difficulty.

The practical implication is direct: for MT evaluation, structured reasoning should not be assumed to improve judgment quality. Direct prompting is both computationally cheaper and empirically more reliable.

10 Conclusion

We have examined gemini 3.1 Flash lite as a judge for English-to-Hindi MT evaluation, finding that LLM-based evaluation is a genuinely useful but imperfect tool whose failures are not where they are commonly assumed to be. The judge is not easily deceived by fluency manipulation; it is more likely to fail by penalizing correct-but-rough translations. Chain-of-thought prompting degrades performance. Rubric-based protocols outperform single-score direct assessment. And perhaps most importantly: human annotation is itself unreliable on challenging inputs, which fundamentally complicates the use of human labels as a gold standard for hard evaluation categories.

For practitioners, these findings support the following recommendations. Use rubric-based LLM evaluation rather than direct scoring when human correlation is the target. Do not assume CoT improves evaluation quality for translation tasks. Always report exact prompt wording, since persona framing can shift scores by up to 12% of the full scale. And invest in understanding human annotation variability before attributing LLM, human gaps solely to judge failure.

Limitations

The study uses a single judge model and a single language pair. Findings regarding judge robustness may not generalize to other models or to language pairs with less world-knowledge coverage. The adversarial manipulations are relatively coarse; subtler semantic perturbations are an important direction for future work. Human annotation was conducted with three annotators; higher annotator counts would yield more stable agreement estimates.

References

Tom Kocmi and Christian Federmann. 2023. Large language models are state-of-the-art evaluators of translation quality. In *Proceedings of EAMT*.

Klaus Krippendorff. 2011. Computing Krippendorff's alpha-reliability. *Departmental Papers (ASC)*.

J Richard Landis and Gary G Koch. 1977. The measurement of observer agreement for categorical data. *Biometrics*, 33(1):159–174.

Yang Liu, Dan Iter, Yichong Xu, Ruochen Xu, and Chenguang Zhu. 2023. G-Eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of EMNLP*, pages 2511–2522.

NLLB Team, Marta R Costa-jussa, James Cross, and 1 others. 2022. No language left behind: Scaling human-centered machine translation. *arXiv preprint arXiv:2207.04672*.

OpenAI. 2025. GPT open source model series. Technical report, OpenAI.

Qwen Team. 2025. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*.

Chenhui Shen and 1 others. 2023. Large language models are not yet human-level evaluators for abstractive summarization. *arXiv preprint arXiv:2305.14562*.

Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, and 1 others. 2023. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, volume 36.

A Annotation Rubric

The following rubric was distributed to all three annotators before annotation.

Fluency (1, 5).

5 Perfectly natural; indistinguishable from native speaker output.
4 Minor awkwardness; meaning is clear throughout.
3 Some unnatural phrasing; understandable but noticeably non-native.
2 Significant grammatical errors; reading requires effort.
1 Ungrammatical or incoherent; meaning cannot be recovered.

Adequacy (1, 5).

5 All meaning preserved; nothing added or omitted.
4 Nearly complete; only trivial details differ.
3 Most meaning present; some information lost or added.
2 Significant meaning distorted or absent.
1 Translation does not correspond to the source.

Note for noisy inputs. Evaluate based on the *intended* clean meaning of the source, not the noisy surface form. A translation that correctly conveys the intended meaning should score highly even when the input was garbled.

B Judge Prompts

All prompts include the English source and Hindi translation. Reference is included except in the no-reference sensitivity variant. Temperature is fixed at 0.1.

Direct scoring.

```
You are an expert evaluator of
English-to-Hindi machine translation.
Source (English): {source}
Hindi Translation: {translation}
Reference (Hindi): {reference}
Rate this translation 1, 5: 1=Very
poor, 2=Poor, 3=Acceptable, 4=Good,
5=Excellent.
Respond with a SINGLE INTEGER only.
```

Pairwise comparison.

```
You are an expert evaluator of
English-to-Hindi machine translation.
Source (English): {source}
Reference (Hindi): {reference}
Translation A: {trans_a}
Translation B: {trans_b}
Which is better overall (accuracy +
fluency)?
Respond with ONE word only: A, B, or
Tie.
```

Rubric-based.

```
Score each criterion 1, 5: adequacy
(meaning preserved?), fluency
(natural Hindi?), terminology
(names/numbers/idioms correct?),
style (register matches source?).
Respond ONLY in this JSON: {"adequacy":X,
"fluency":X, "terminology":X,
"style":X}
```

Chain-of-thought.

```
Think step by step:
Step 1: Is the full meaning of the source
preserved?
Step 2: Are names, numbers, dates, and
key terms correct?
Step 3: Is the Hindi grammatically
correct and natural?
Step 4: Does it match the register of
the source?
Step 5: Assign a final score from 1 to
5.
Format: REASONING: <analysis> SCORE:
<integer>
```

C Dataset Category Examples

Ambiguous. "Her first was the Slalom, where she earned a Did Not Finish in her first run. 36 of the 116 competitors had the same result in that race." The numerical reference creates ambiguity about which race is being described; translators must commit to a reading that may not be uniquely supported by the source.

Hallucination-prone. "The original population hasn't changed at all, they still need the same adaptations as before." The pronoun "they" creates reference ambiguity; both systems must choose a grammatical number and gender that is not specified in English.

Noisy input (OCR type). *Original:* "They all ran back from where the accident had occurred." *Noisy:* "They a11 ran back fr0m where theaccident had occurred." Character substitutions ($l \to 1$, $o \to 0$) and a merged word test system robustness to OCR-style corruption.

Negation/quantifier. "No matter how docile they may look, bison, elk, and moose are wild animals." The concessive construction ("no matter how") must be preserved; dropping or softening it changes the pragmatic force of the sentence.

Long and complex. "Four skiers in the women's sitting group failed to finish their runs. 36 of the 116 competitors had the same result in that race." Two sentences, two embedded numerical facts, and a negation ("failed to finish"), all of which must be preserved faithfully in a language with different word order.

D Adversarial Test Suite

Type	English Source	Manipulation
Fluent-wrong	The patient was not allowed to leave the hospital.	"leave" → "go to" [direction]
Fluent-wrong	The population increased by ten percent.	"increased" → "decreased"
Fluent-wrong	She received the highest award in the competition.	"received" → "gave"
Fluent-wrong	The treaty was signed before the war ended.	"before" → "after"
Correct-poor	Scientists discovered a new species of fish.	Correct meaning; heavy English borrowing
Correct-poor	The government announced new economic policies.	Correct meaning; unnatural passive
Correct-poor	Children played in the park.	Correct meaning; wrong plural form
Correct-poor	He completed his studies in three years.	Correct meaning; scrambled order
Meaning shift	Almost all students passed the exam.	"almost all" → "all"
Meaning shift	The medicine should be taken after meals.	"after" → "before"
Meaning shift	The suspect could have committed the crime.	"could have" dropped [modal]
Meaning shift	He rarely misses his morning workout.	"rarely" → "sometimes"

Table 5: Full adversarial test suite. The judge correctly assigned higher scores to the reference translation in all 12 cases.